

ECE-GY 9143 - High Performance Machine Learning

Homework Assignment 4

Zehra Sura and Robert Kingan

Due Date: Nov 7, 2025

Fall 2025

Max Points: 100

Instructions

This lab is intended to be performed individually; great care will be taken in verifying that students are authors of their own submissions.

This lab has three parts: Part-A on vector addition and matrix multiplication in CUDA, Part-B on CUDA Unified Memory, and Part-C on convolution in CUDA.

Submit a zip file with all the code files and the report. The format for the zip file is [netid].zip, with underlying folders for each part. The code file for individual questions should be in the format q[number].[ext].

Part-A: CUDA Matrix Operations (40 points)

The purpose of this exercise is for you to learn how to write programs using the CUDA programming interface, how to run such programs using an NVIDIA graphics processor, and how to think about the factors that govern the performance of programs running within the CUDA environment. For this assignment, you will modify two provided CUDA kernels. The first is part of a program that performs vector addition. The second is part of a program to perform matrix multiplication. The provided vector addition program does not coalesce memory accesses. You will modify it to coalesce memory access. You will modify the matrix multiplication program to investigate its performance and how it can be optimized by changing how the task is parallelized. You will create an improved version of the matrix multiplication program and empirically test the time it takes to run. You will analyze the results of your timing tests.

You will turn in all code you have written and used for this assignment, a makefile that will compile your code, and a standard lab report documenting your experiments and results. We need to be able to compile your code by executing "make".

Running on HPC

Upload the PartA.zip file to your scratch directory and unzip it. In order to compile, uncomment the lines corresponding to the required program in the makefile. For example, in order to run vecadd00, uncomment these 4 lines in the makefile, and ensure there is a tab indent for the NVCC lines

```
vecaddKernel100.o: vecaddKernel100.cu
    ${NVCC} $< -c -o $@ $(OPTIONS)

vecadd00: vecadd.cu vecaddKernel.h vecaddKernel100.o timer.o
    ${NVCC} $< vecaddKernel100.o -o $@ $(LIB) timer.o $(OPTIONS)
```

You can use these commands in the batch file:

```
cd path/to/PartA
```

```
make clean  
make  
./vecadd00 arguments  
./matmul00 arguments
```

Problem 1: Vector Add and Coalescing Memory Access

You are provided with vecadd, a micro benchmark to determine the effectiveness of coalescing. The provided code is a non-coalescing version, and it is your job to create a coalescing version, and to measure the difference in performance of the two codes.

Here is a set of files provided in PartA.zip for this problem

- Makefile
- vecadd.cu
- vecaddKernel.h
- vecaddKernel00.cu
- timer.h
- timer.cu

Q1

5 points

Compile and run the provided program as `vecadd00` (using the kernel `vecaddKernel00`) and collect data on the time the program takes with the following number of values per thread: 500, 1000, 2000.

Include a short comment in your `Makefile` describing what your various programs do, and fully describe and analyse the experiments you performed in your lab report.

Q2

10 points

In `vecadd01` you will use a new device kernel, called `vecAddKernel01.cu`, to perform vector addition using coalesced memory reads. Change the `Makefile` to compile a program called `vecadd01` using your kernel rather than the provided kernel. Modify the `Makefile` as appropriate so the `clean` and `tar` commands will deal with any files you have added.

Test your new program and compare the time it takes to perform the same work performed by the original. Note and analyse your results, and document your observations in the report.

Problem 2: Shared CUDA Matrix Multiply

For the next part of this assignment, you will use matmultKernel00 as a basis for another kernel with which you will investigate the performance of matrix multiplication using a GPU and the CUDA programming environment. Your new kernel should be called matmultKernel01.cu. Your code files must include internal documentation on the nature of each program, and your README file must include notes on the nature of the experiment. You can have intermediate stages of this new kernel, eg without and with unrolling, but we will grade you on your final improved version of the code and its documentation.

Here is a set of provided files in PartA.zip for this problem:

- Makefile
- matmult.cu
- matmultKernel.h
- matmultKernel00.cu
- timer.h
- timer.cu

In both the questions of this problem (Q3 and Q4 below) you should investigate how each of the following factors influences performance of matrix multiplication in CUDA:

1. The size of matrices to be multiplied
2. The size of the block computed by each thread block
3. Any other performance enhancing modifications you find

Q3

8 points

First you should time the initial code (with the provided kernel `matmultKernel00`) using square matrices of each of the following sizes: 256, 512, 1024. When run with a single parameter, the provided code multiplies that parameter by `FOOTPRINT_SIZE` (set to 16 in `matmult00`) and creates square matrices of the resulting size. This was done to avoid nasty padding issues: you always have data blocks perfectly fitting the grid.

Q4

12 points

In your new kernel each thread computes four values in the resulting C block rather than one. So now the `FOOTPRINT_SIZE` becomes 32. (Notice that this is taken care of by the `Makefile`.) You will time the execution of this new program with matrices of the sizes listed above to document how your changes affect the performance of the program. Provide the speedup that your new kernel achieves over the provided kernel.

To get good performance, you will need to be sure that the new program coalesces its reads and writes from global memory. You will also need to unroll any loops that you might be inclined to insert into your code.

Q5

5 points

Can you formulate any rules of thumb for obtaining good performance when using CUDA? Answer this question in your conclusions.

Part-B: CUDA Unified Memory (20 points)

In this problem we will compare vector operations executed on host vs on GPU to quantify the speed-up.

Q1

4 points

Write a C++ program that adds the elements of two arrays with a K million elements each. Here K is a parameter of the program. Profile and get the time to execute this program for K=1,5,10,50,100. Use `free()` to free the memory at the end of the program.

Q2

6 points

Next using CUDA execute the add operation as a kernel on GPU. Use `cudaMalloc()` to allocate memory on GPU for storing the arrays and `cudaMemcpy()` to copy data to and

from the GPUs. Note that host and GPU memory is not shared in this case. You will be running three scenarios for this part:

1. Using one block with 1 thread
2. Using one block with 256 threads
3. Using multiple blocks with 256 threads per block, with the total number of threads across all blocks equal to the size of arrays.

Again, profile and get the time to execute this program for $K = 1, 5, 10, 50, 100$. So for each of the three scenarios above you will get profile times for five different values of K . Use `cudaFree()` and `free()` to free the memory on the device and host at the end of the program.

Q3 6 points

This time you will repeat Step 2 using CUDA Unified Memory and providing a single memory space accessible by all GPUs and CPUs in your system. Instead of `cudaMalloc()` you will use `cudaMallocManaged()` to allocate data in unified memory, which returns a pointer that you can access from host (CPU) code or device (GPU) code. To free the data, just pass the pointer to `cudaFree()`. Again, profile and get the time to execute this program for $K = 1, 5, 10, 50, 100$ for each of the three scenarios.

Q4 4 points

Plot two charts — one for Step 2 (without Unified Memory) and one for Step 3 (with Unified Memory). The x-axis is the value of K (in millions) and the y-axis is the time to execute the program (one chart for each of the three scenarios). In both charts also plot the time to execute on CPU only. Since the execution time scale on CPU and GPU may be orders of magnitude different, you may want to use log-log scale for the y-axis when plotting.

For this problem you may want to refer to slide deck for **Lecture 6** which has a similar example and the following blogpost: An Even Easier Introduction to CUDA, available at: <https://developer.nvidia.com/blog/even-easier-introduction-cuda/>.

Part-C: Convolution in CUDA (40 points)

C1: Convolution in CUDA (15 points)

In this part of the lab we implement a convolution of an image using a set of filters. Consider the following:

- An input tensor I with dimensions: C, H, W . Each element of I is generated as follows:

$$I[c, x, y] = c(x + y)$$

- A set of convolution filters with dimensions: K, C, FH, FW . Each element of the filter F is generated as follows: $F[k, c, i, j] = (c + k)(i + j)$
- Dimensions are: $H = W = 1024, C = 3, FW = FH = 3, K = 64$
- All tensors have double elements (double precision)

- A tensor I_0 of sizes $C, W+2P, H+2P$ where $P=1$. I_0 is obtained from the tensor I adding the padding rows and columns with the elements set to zero.
- The output tensor O with dimensions: K, W, H . Each pixel of the output tensor

$$O[k, x, y]$$

is obtained as:

$$O[k, x, y] = \sum_{c=0}^{C-1} \sum_{j=0}^{FH-1} \sum_{i=0}^{FW-1} F[k, c, FW - 1 - i, FH - 1 - j] I_0[c, x + i, y + j]$$

Note that we need to use the transpose of the filter in order to compute a convolution rather than a cross-correlation. Implement a simple convolution algorithm in CUDA without tiling and without shared memory.

- Print the checksum as the total sum of the elements of O along all its dimensions.
- Report the time to execute the CUDA kernel with the convolution, without including the time to generate the data and to copy it into CUDA memory (only the kernel execution time – you need to use a synchronization primitive).

C2: Tiled Convolution with CUDA (15 points)

Implement the convolution at exercise C1 using shared memory and tiling.

- Print the checksum as the sum of the elements of O along all its dimensions. The checksum is expected to be the same as C1.
- Report the time to execute the CUDA kernel with the convolution, without including the time to generate the data and to copy it into CUDA memory (only the kernel execution time – you need to use a thread synchronization primitive).

C3: Convolution with cuDNN (10 points)

Implement convolution at exercise C1 using cuDNN with the CUDNN_CONVOLUTION_FWD_PREFER_FASTEST algorithm selection.

- Print the checksum as the sum of the elements
- Report the time to execute the CUDA kernel with the convolution, without including the time to generate the data and to copy it into CUDA memory (only the kernel execution time).

Program Output Format

A single binary generated from a single .cu file (no headers) should be used for the 3 exercises. The program output should look as follows:

```
C1_checksum,C1_execution_time
C2_checksum,C2_execution_time
C3_checksum,C3_execution_time
```

Where each string contains the checksum and the execution time of the exercises C1, C2 and C3 respectively. Execution times should be printed in milliseconds with 3 decimals. Failing to respect this format is -3 points.